

DATA WRANGLING PROJECT

ANDRÁS SOMI, 2017. JULY - UPDATED VERSION 2

Data

- `budapest_hungary_inner.osm` (107 MB)
- `dump.json` (114 MB)

I created a custom Metro extract on Mapzen for the inner parts of Budapest, capitol of Hungary. The available basic extract of Budapest area was way bigger than this exercise requires (close to 1GB) and also contained areas that in reality are not part of the city (that may cause conflicts in postcodes or street name duplications).

I live in the city, I have a general understanding of the naming conventions, special characters, etc. Also the different language makes it difficult to simply reuse code snippets from the examples, so I need to thoroughly think through every piece of it.

Source: https://mapzen.com/data/metro-extracts/metro/budapest_hungary/

Auditing data

First, to have a general understanding of the data I ran through the whole xml counting the occurrences of different attributes by tags and the overall count of the different tags. The dataset seems to be pretty uniform, ie. the different type of tags tend to have the same set of attributes for all the instances. For node tags `uid` and `username` attributes are missing in a tiny fraction of the cases.

TAG AND ATTRIBUTE COUNTS:

```
{'bounds': {'attributes': {'maxlat': 1, 'maxlon': 1, 'minlat': 1,
'count': 1},
'member': {'attributes': {'ref': 68919, 'role': 68919, 'type':
68919},
'count': 68919},
'nd': {'attributes': {'ref': 547248}, 'count': 547248},
'node': {'attributes': {'changeset': 422629,
'id': 422629,
```

```

        'lat': 422629,
        'lon': 422629,
        'timestamp': 422629,
        'uid': 422611,
        'user': 422611,
        'version': 422629},
    'count': 422629},
    'osm': {'attributes': {'generator': 1, 'timestamp': 1, 'version':
1},
        'count': 1},
    'relation': {'attributes': {'changeset': 7098,
        'id': 7098,
        'timestamp': 7098,
        'uid': 7098,
        'user': 7098,
        'version': 7098},
        'count': 7098},
    'tag': {'attributes': {'k': 402982, 'v': 402982}, 'count':
402982},
    'way': {'attributes': {'changeset': 75148,
        'id': 75148,
        'timestamp': 75148,
        'uid': 75148,
        'user': 75148,
        'version': 75148},
        'count': 75148}}

```

But most of the interesting data is stored in `tag` tags as key-value pairs in the `k` and `v` attributes. Let's look into these.

Auditing street names

For a 'gold standard' of types of public places I used the information from [Wikipedia](#). This seems to be a complete, official list of the Hungarian names for different types of streets, roads and other public areas.

It's a pretty long list but most of the datapoints should end with one of the three most popular types (street, road, square). (In Hungarian we just put the type after the name of the street in small caps like 'Ilka *utca*' or 'Döbrentei *tér*')

I found 18 names that don't fit into any of the official categories, but most of them actually make sense as they are grammatical variations of basic types (eg. *'útja'* means the road of someone or something, so absolutely valid). The rest are some unique places like a castle (yes, there are castles in Budapest!) with their unique names.

UNEXPECTED STREET NAMES:

```
{'Erzsébet királyné útja',
 'Ferenciek tere',
 'Hadak útja',
 'Harminckettesek tere',
 'Ifjúság útja',
 'Kucsma', # THIS IS INCORRECT
 'Kunigunda útja',
 'Magyar tudósok körútja',
 'Margitsziget',
 'Nagy Lajos király útja',
 'Népliget',
 'Népliget aluljáró',
 'Rákos-patak',
 'Rákospalotai körvasútsor',
 'Rózsák tere',
 'Vajdahunyadvár',
 'Vigadó téri hajóállomás',
 'Árpád fejedelem útja'}
```

We have only one entry that is suspicious: *'Kucsma'*. That's actually the name of the street not the type, and apparently *'utca'* (street) is missing from the end. *'Kucsma utca'* occurs seven times in the dataset.

The search also brought up a few cases where the street name starts with lower case letter. This should be also handled before uploading the dataset to a database.

```
...
'Zöldmáli lejtő': 7,
'dessewffy utca': 1, # SMALL CAPS
'podmaniczky utca': 1, # SMALL CAPS
```

```
'szabolcs utca': 1, # SMALL CAPS
'százados út': 1, # SMALL CAPS
'Ábel Jenő utca': 4,
...
```

UPDATE: After first review I found a bug that limited the street name and postcode audit to <way> tags. After fixing it some other odd street names popped up (eg. *'Táncsics Mihály utca 5'* and *'Városmajor utca 5. fsz. 3.'*) where the street name field contains the house number too.

Auditing postcodes

In Hungary we use four-digit postcodes. All Budapest postcodes start with 1 and the second and third digit denotes the number of district. There are 23 districts in the city, so the inner two digits should be between 01 and 23, except for the island called Margitsziget, where the inner two digits are 00.

There are no residential areas on the island but there are some cultural and sports facilities so we should still expect a few 00s to pop up in the dataset. ([The Districts of Budapest \(Wikipedia\)](#))

Four odd postcodes popped up in the audit. In the last one 'H' denotes Hungary in an international postcode format. It's a pattern we can easily correct programatically in the dataset.

1503 and 1507 are most likely typos for 1053 and 1057 (confirmed by the street names) while 1476 seems to be a valid postcode even though it does not seem to adhere to the standard format (but Google also gives [valid results](#))

```
{'1476': {'count': 1, 'tags': ['Üllői út']},
'1503': {'count': 1, 'tags': ['Kérő utca']},
'1507': {'count': 1, 'tags': ['Irinyi József utca']},
'H-1026': {'count': 2, 'tags': ['Pasaréti út', 'Pasaréti út']}}
```

UPDATE: Auditing email addresses

Strict email validation can be tricky, but obvious outliers can be filtered out by using a simple regular expression to control for the necessary elements (eg. '@', domain, extension). This surfaced only one entry, which turned out to be valid, my regex just did not include '_' as a valid character.

INVALID EMAIL ADDRESSES:

```
['fovarosi_keptar@mail.btm.hu', 'fovarosi_keptar@mail.btm.hu']
```

UPDATE: Auditing phone numbers

Just by printing out all the phone numbers it's quite clear this data is really a mess. At least a handful of patterns are visible and quite a few odd formats are also present (and some obvious errors, like email addresses stored as phone numbers). Most of this can be transformed programatically but many of the entries have to be cleaned by hand.

Some examples of the phone number field in the dataset:

- +36 1 413 0652
- +36 1 4131310
- '+36 21 3813980,+36 21 3813981,+36 21 3813977'
- '+361-952-0432'
- '+3612001348'
- '+3670 520 99 53',
- '437-9011'
- '23948400242'
- '(06-1) 459-2302'
- '(1) 815 1100'
- '+36-1-213-9039'

I prefer having phone numbers in +36 xx xxx xxxx format, with only whitespaces, no hyphens or brackets. Please note that landlines in Budapest follow the +36 1 xxx xxxx pattern, while mobile numbers and landlines outside Budapest are in a +36 xx xxx xxxx format (one extra digit after the country code). So the first one is correct above, but not the others.

In `cleaning.py` I managed to programatically format most of the phone numbers that otherwise seemed correct, but a handful of cases still remain (eg. several phone numbers in a single field, missing or extra digits, etc.).

Auditing coordinates

I audited latitude and longitude coordinates to be float numbers around 47.5 and 19 respectively. Not surprisingly no odd coordinates popped up thanks to the way the data was obtained.

Cleaning the data

In `cleaning.py` I performed several cleaning steps during the shaping of the OSM XML elements into Python dictionaries before appending them to the output JSON file. I give a brief description here, but the commented code is also self-explanatory.

1. Cleaning streetnames

The main issue was lowercase streetnames, which can be detected by checking whether the streetname equals its own lowercase variant (meaning there's no uppercase letter in it). Then the transformation is made by creating a new string from the uppercase initial letter and the rest. I also replaced two instances where *'utca'* (street) was missing from the end.

```
def clean_streetname(streetname):
    '''Programatically fix issues of street names'''
    if streetname in ['Kucsma', 'Dohány']:
        new_streetname = streetname + ' utca'
    elif streetname.lower() == streetname:
        new_streetname = streetname[0].upper() + streetname[1:]
    else:
        new_streetname = streetname
    return new_streetname
```

2. Cleaning postcodes

Although not many problematic postcodes reared their head during the audit, there was one specific pattern, the appearance of *'H-'* at the beginning, that should be programatically corrected by returning a slice of the string. I also replaced two instances of wrong postcodes (probably typos).

```
def clean_postcode(postcode):
    '''Programatically fix issues with postcodes'''
    postcode_string = str(postcode)
    if postcode_string[:2] == 'H-':
        return int(postcode_string[2:])
    elif postcode_string == '1503':
        return 1053
    elif postcode_string == '1507':
        return 1057
```

```
else:
    return int(postcode)
```

3. *Cleaning phone numbers*

That was the toughest bit, as lots of different patterns appeared in the dataset. These were the steps to standardize them into the preferred (either `+36 xx xxx xxxx` or `+36 1 xxx xxxx` format).

1. Define a regular expression for the preferred pattern.
2. Remove special characters (`/`, `(`, `)`, `-`) and all the whitespaces.
3. Replace at the beginning the `'0036'`, `'036'`, `'006'` or `'06'` digits of country code with the preferred format `('+36')`.
4. Or add `'+36'` to the beginning if it's still missing (because the original string did not contain any form of country code, so there were nothing to replace in the previous step).
5. We now have a sort of (semi-)standardized format from which we can pick those strings that only contain the proper number of digits (8 or 9) besides the `'+36'` at the beginning.
6. Add some whitespace to these to get the preferred format and return the formatted string.
7. Log an error for the rest that cannot be standardized this way.

This process handles the majority of the erroneous phone numbers, though some patterns still can be found in the rest (eg. multiple phone numbers in one row or an extension added with some special characters) that might also be handled in a programatic way.

```
def clean_phone_numbers(phone_number):
    '''Transform phone number to +36 1 xxx xxxx or +36 xx xxx xxxx
    format.'''
    preferred_format = '\\+36\\s[1-9]0?\\s[0-9]{3}\\s[0-9]{4}$'
    match = re.match(preferred_format, phone_number)

    # Return if the phone number already follows the preferred
    pattern
    if bool(match):
        return phone_number

    # Remove special characters and whitespaces
    stripped = re.sub('[/()-]', '', phone_number).replace(' ', '')
```

```

# Replace the country code with the right format
replaced = re.sub('^0036|^06|^006|^036', '+36', stripped)

# Insert country code to the beginnig if it's missing
# but the number otherwise seems good (strictly 8 or 9 digits)
if replaced[:3] != '+36' and re.match('^[0-9]{8,9}$',
replaced):
    replaced = '+36' + replaced

# Budapest landlines
if len(replaced) == 11 and replaced[:3] == '+36' and
replaced[3:4] == '1':
    formatted = replaced[:3] + ' ' + replaced[3:4] + ' ' +
replaced[4:7] + ' ' + replaced[7:]

# Mobile or non-Budapest landline
elif len(replaced) == 12 and replaced[:3] == '+36':
    formatted = replaced[:3] + ' ' + replaced[3:5] + ' ' +
replaced[5:8] + ' ' + replaced[8:]

# If it doesn't fit into any categories log the error and
return the original
else:
    logging.error('Cannot process phone number:
{0}'.format(phone_number))
    return phone_number

return formatted

```

Querying the data

Number of documents

```

> db.budapest.find().count()
504875

```

Count by type of tag


```

> db.budapest.aggregate([
  {'$group': {
    '_id': '$type',
    'count': {'$sum': 1}
  }},
  {'$project': {
    '_id': 0,
    'type': '$_id',
    'count': 1
  }}
])

[{'count': 7098, 'type': 'relation'},
 {'count': 75148, 'type': 'way'},
 {'count': 422629, 'type': 'node'}]

```

Top 5 postcodes by number of occurrence

The 11th district is one of the biggest one (maybe the biggest one), no surprise that x11x postcodes are the most frequent in the dataset.

```

> db.budapest.aggregate([
  {'$match': {
    'address.postcode': {'$exists': True}
  }},
  {'$group': {
    '_id': '$address.postcode',
    'count': {'$sum': 1}
  }},
  {'$sort': {'count': -1}},
  {'$limit': 5},
  {'$project': {
    '_id': 0,
    'postcode': '$_id',
    'count': 1
  }}
])

[{'count': 1922, 'postcode': 1112},
 {'count': 1566, 'postcode': 1118},

```

```
{'count': 1035, 'postcode': 1124},
{'count': 1025, 'postcode': 1025},
{'count': 735, 'postcode': 1089}]
```

Number of contributing users

Not too many compared to the population of 2 million people... python

```
> db.budapest.distinct('created.user').length
1232
```

Top 3 users

igor2 is the big winner with 3 times more entries than the second most active user.

```
> db.budapest.aggregate([
  {'$group': {
    '_id': '$created.user',
    'count': {'$sum': 1}
  }},
  {'$sort': {'count': -1}},
  {'$limit': 3},
  {'$project': {
    '_id': 0,
    'username': '$_id',
    'count': 1
  }}
])

[{'count': 134944, 'username': 'igor2'},
 {'count': 33466, 'username': 'MartinHun'},
 {'count': 23452, 'username': 'vasony'}]
```

Number and type of amenities

Apparently Budapest is the city of benches.

```
> db.budapest.find({'amenity': {'$exists': True}}).count()
10271

> db.budapest.aggregate([
  {'$match': {
```

```

        'amenity': {'$exists': True}
    }},
    {'$group': {
        '_id': '$amenity',
        'count': {'$sum': 1}
    }},
    {'$sort': {
        'count': -1
    }},
    {'$limit': 5},
    {'$project': {
        '_id': 0,
        'type': '$_id',
        'count': 1
    }}
  ])

[{'count': 1394, 'type': 'bench'},
 {'count': 975, 'type': 'restaurant'},
 {'count': 893, 'type': 'bicycle_parking'},
 {'count': 791, 'type': 'parking'},
 {'count': 695, 'type': 'waste_basket'}]
```

Most popular cuisines

Italian is our choice after the local tastes. (I guess pizza delivery is still a big business here.)

```

db.budapest.aggregate([
    {'$match': {
        'amenity': 'restaurant',
        'cuisine': {'$exists': True}
    }},
    {'$group': {
        '_id': '$cuisine',
        'count': {'$sum': 1}
    }},
    {'$sort': {'count': -1}},
    {'$limit': 3},
    {'$project': {
        '_id': 0,

```

```

        'cuisine': '$_id',
        'count': 1
    }}
])

[{'count': 76, 'cuisine': 'regional'},
 {'count': 61, 'cuisine': 'hungarian'},
 {'count': 54, 'cuisine': 'italian'}]
```

Further ideas

Dates and timestamps

It would be useful to have the timestamps in date format in MongoDB, but Python `datetime` objects cannot be serialized to JSON so in this workflow of importing the data at once into MongoDB we shouldn't change the timestamps from strings to dates.

```

entry = db.budapest.find_one({'created.timestamp': {'$exists':
True}})
timestamp = entry['created']['timestamp']
print(isinstance(timestamp, str))
True
```

There might be several ways to handle this:

1. Before creating the JSON file transform string timestamps to UNIX timestamps and store them as integers. This makes somewhat easier to create date-based queries and still doesn't break the JSON dump (still not the most convenient way to handle timestamps).
2. After importing the JSON data to MongoDB run a script that transforms the string timestamps to proper `datetime` objects and updates the appropriate field document-by-document. This would be a time-consuming operation but then we can use the timestamps as dates in queries and aggregations.

UPDATE: Improving the data with public transport lines

Even though the dataset contains 1414 nodes tagged as bus-stops, the numbering of available bus lines are missing on these nodes. The dataset also contains 7 bigger bus stations, but only one of them has information on which bus lines are

available there. It would be useful to enrich the dataset with bus, tram and local train line information.

Benefits: * The data could be used for public transport mapping and travel planning applications. * It should be available at the local transportation company (BKK) in a complete and high quality format (but maybe not free), so no need for 'crowdsourcing' the initial upload.

Possible problems: * The data might change frequently (either temporarily or permanently) so frequent updates and monitoring is needed (that's where crowdsourcing and gamification might come into play) * The public transport system of Budapest is huge which might produce vast amount of data to process.

```
> db.budapest.find({'highway': 'bus_stop'}).count()
1414

> db.budapest.find({'amenity': 'bus_station', 'lines': {'$exists':
True}})
[{'_id': ObjectId('5975a8e22bfafbe107b8375e'),
  'amenity': 'bus_station',
  'created': {'changeset': '44636379',
              'timestamp': '2016-12-24T00:04:59Z',
              'uid': '2982309',
              'user': 'vasony',
              'version': '7'},
  'id': '287438376',
  'lines': '8, 139, 153, 153A, 154, 239',
  'name': 'Gazdagréti lakótelep',
  'operator': 'BKV',
  'pos': [47.473576, 18.9926846],
  'type': 'node',
  'wheelchair': 'yes'}]
```